

Lab 01 – Computer Vision

BSDSF21M020 – BUSHRA SHAHBAZ

1st October, 2024

Task#1: Run `1draw.py`. State your observations.



1. **Blank Image:** A black canvas of size 500x500 pixels is created using `np.zeros()`. This blank image is shown in a window labeled "Blank," with all pixel values initially set to zero (black).
2. **Red Square:** A section of the image, from coordinates (200, 300) to (300, 400), is filled with red ((0, 0, 255) in BGR format). This small red square is added to the canvas, and the updated image is shown in a window labeled "Green."
3. **Green Rectangle:** A larger green rectangle is drawn, starting at the top-left corner (0, 0) and covering half of the canvas (the top-left quadrant). This green rectangle ((0, 255, 0)) is filled in completely due to the thickness being set to -1.
4. **Red Circle:** A red circle is drawn at the center of the canvas, with a radius of 40 pixels. It overlaps the green rectangle because both are placed at the center of the image.
5. **White Line:** A white line is drawn diagonally across the canvas, starting from coordinates (100, 250) and ending at (300, 400), with a thickness of 3 pixels. The line cuts across the existing shapes on the image.
6. **Text:** Green text, "Hello, my name is Computer Vision!!!", is added near the center of the image. It uses the `FONT_HERSHEY_TRIPLEX` font and is positioned at coordinates (0, 225).

Each element is drawn on the same canvas, and the updates are displayed progressively using `cv.imshow()`. The `waitKey(0)` function keeps the windows open until a key is pressed.

Task #2: Run 2imageSize.py. State your observations.

1. **Opening the Image:** The script opens the image file "cat_hat.jpg" using `Image.open()`. The image is loaded and stored in the variable `image`.
2. **Image Dimensions:** The size of the image is retrieved using the `image.size` attribute. This returns a 2-tuple containing the width and height of the image in pixels.
3. **Displaying Size and Dimensions:** The width and height are then printed individually using `print(width, height)`, and the full size (tuple) is also printed via `print(image.size)`.

In summary, the script loads the image, retrieves its dimensions, and prints them. This demonstrates how Pillow allows for easy manipulation and access to image properties.

```
In [2]: """
Created on Fri Sep 27 22:59:36 2024

@author: HP
"""

from PIL import Image

filename = "cat_hat.jpg"
with Image.open(filename) as image:
    width, height = image.size
    print(image.size)
    print(width,height)
#Image.size gives 2-tuple and the width, height can be obtained

(300, 168)
300 168
```

Task #3: Run 3MyFirst.py. State your observations.

- **Prints a message:** "My First Python Program".
- **Creates a NumPy array:** A 3x4 matrix using `np.arange(12).reshape(3,4)`.
- **Displays the array as an image:** Using `plt.imshow()` with a grayscale color map.
- **Removes axes:** Using `plt.axis('off')`.
- **Shows the image:** With `plt.show()`, visualizing the matrix as a grayscale image where pixel intensity reflects array values.

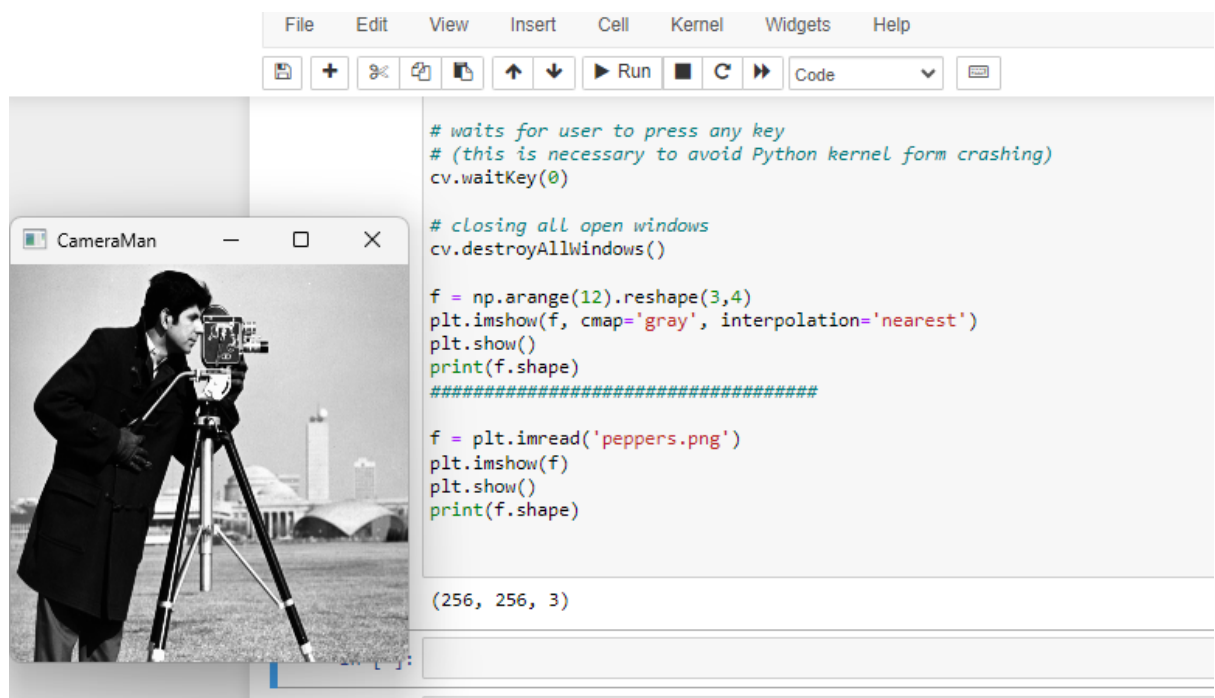
```
plt.axis('off')  
plt.show()
```

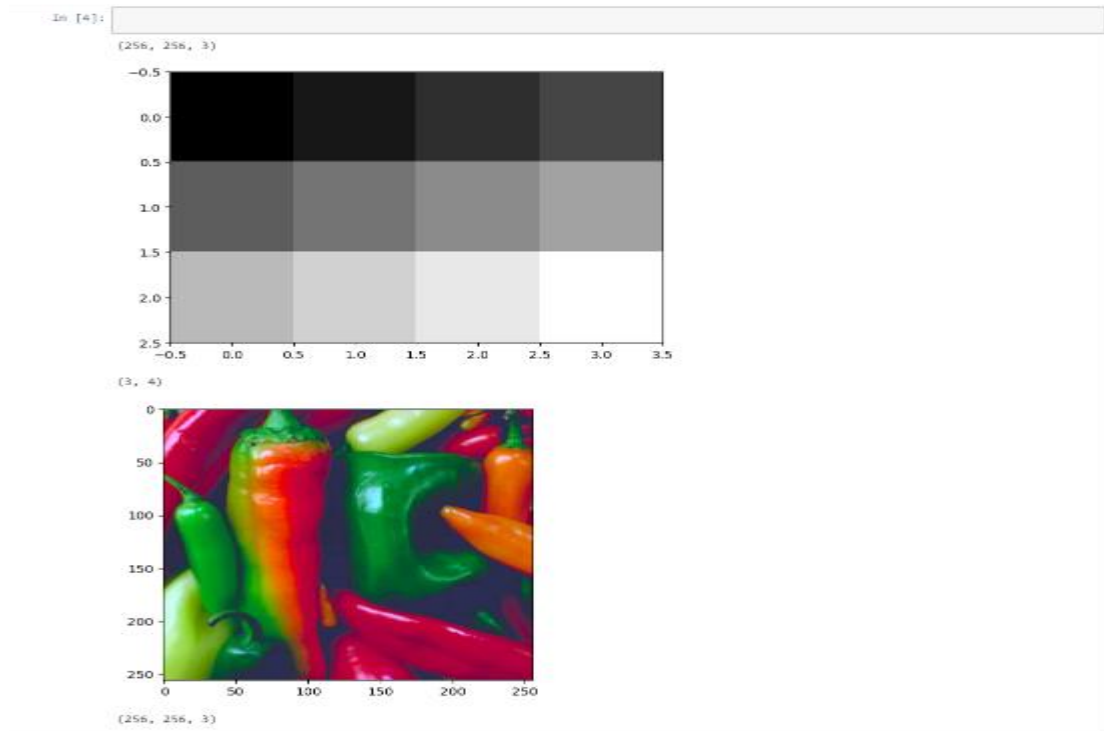
```
My First Python Program  
[[ 0  1  2  3]  
 [ 4  5  6  7]  
 [ 8  9 10 11]]
```



Task #4: Run 4Read_write.py. State your observations.

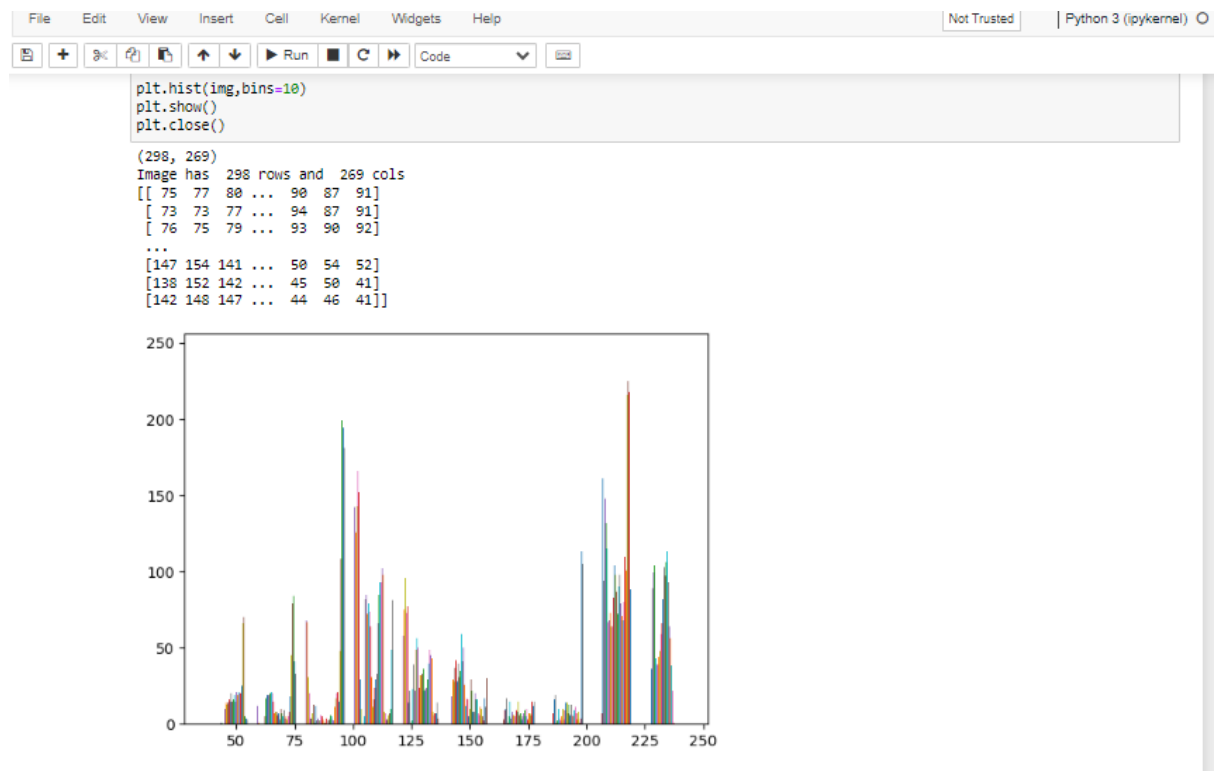
- ☐ **Reads and Displays an Image:**
 - The image "cameraman.png" is loaded using `cv.imread()`.
 - It is displayed in a window named 'CameraMan' using `cv.imshow()`.
 - The dimensions of the image (height, width, channels) are printed using `image.shape`.
 - The script waits for a key press with `cv.waitKey(0)` and then closes all windows with `cv.destroyAllWindows()`.
- ☐ **NumPy Array Visualization:**
 - A 3x4 matrix `f` is created using `np.arange(12).reshape(3,4)` and visualized as a grayscale image using `plt.imshow()`. The matrix dimensions are printed via `f.shape`.
- ☐ **Reads and Displays Another Image:**
 - The image "peppers.png" is loaded using `plt.imread()` and displayed using `plt.imshow()`.
 - The dimensions of the image are printed using `f.shape`.





Task #5: Run 5row_col.py. State your observations.

- ☐ **Reads an Image:**
 - The image file 'gray_output.jpg' is read using `io.imread()`, and its shape (dimensions) is printed. This shape indicates the number of rows and columns in the image.
- ☐ **Extracts Dimensions:**
 - The number of rows (`r`) and columns (`c`) are extracted from `img.shape`, where `r` represents the height and `c` represents the width of the image.
 - It then prints the dimensions, indicating how many rows and columns the image has.
- ☐ **Prints Image Values:**
 - The script prints all pixel values of the image stored in the `img` array. This will display the grayscale intensity values if the image is in grayscale.
- ☐ **Histogram of Pixel Values:**
 - A histogram of the pixel values is created using `plt.hist(img, bins=10)`, which shows the distribution of pixel intensities across the image.
 - Finally, `plt.show()` displays the histogram, and `plt.close()` closes the plotting window.



Task #6: Run `6run_tell.py`. State your observations.

1. **Imports Libraries:** It imports `matplotlib.pyplot` for plotting and `skimage.data` for sample images.
2. **Loads and Converts Image:** It loads a sample RGB image of a cat and converts it to grayscale using `rgb2gray`.
3. **Creates Subplots:** It sets up a figure with two subplots to display the original RGB image and the grayscale image side by side.
4. **Displays Images:** Each image is shown with appropriate titles ("RGB" for the original and "Grayscale" for the converted image).
5. **Shows the Plot:** The layout is adjusted for clarity, and the final plot is displayed using `plt.show()`.

